

Artificial Intelligence and Machine Learning 6CS012

Deep Learning Based News Text Classification using RNN, LSTM, and Word2Vec Embeddings

Canvas ID: 2406775
Name: Gaurab Rawat
Group: L5CG18
Lecturer: Durga Pokharel
Tutor: Shiv Kumar Yadav
Submitted on: 2026-05-1

ABSTRACT

This project presents a deep learning approach to automatic news headline classification using Natural Language Processing (NLP) techniques. The primary objective was to develop and compare multiple recurrent neural network architectures for predicting news categories from short textual inputs. Three distinct models were implemented: a Simple RNN with trainable embeddings, an LSTM with trainable embeddings, and an LSTM integrated with pre-trained GloVe-50 embeddings.

The dataset consisted of 11,025 categorized news headlines across five classes: SPORTS, TECH, MONEY, CULTURE & ARTS, and EDUCATION — sourced from the HuffPost News Category Dataset (Variant 4 Slate). Prior to model training, a comprehensive preprocessing pipeline was applied, including lowercasing, URL and mention removal, contraction expansion, stop word removal, and lemmatization. The cleaned text was subsequently tokenized and padded to a fixed sequence length of 12 tokens for deep learning input.

The comparative analysis revealed significant performance differences between architectures. The Simple RNN achieved a test accuracy of 56.64%, struggling with vanishing gradients even over these short sequences. The LSTM improved substantially to 79.46%, with its gating mechanism providing better class separation — particularly recovering MONEY F1 from near-zero. The LSTM combined with frozen GloVe-50 embeddings achieved the highest overall accuracy of 82.49%, demonstrating that representation quality outweighs architecture complexity for short-text tasks. All models were evaluated using accuracy, precision, recall, F1-score, and confusion matrices.

The project further demonstrates a live prediction interface built with Grado, allowing users to classify headlines interactively across all three model pipelines. Key findings confirm that pre-trained embeddings offer the most impactful performance improvement for small-dataset, short-text classification.

Table of Contents

ABSTRACT	2
Table of Figure	4
1. INTRODUCTION	6
2. DATASET	7
2.1 Dataset Source	7
2.2 Dataset Characteristics	7
2.3 Data Preprocessing	8
Lowercasing	8
URL and Mention Removal	8
Number and Special Character Removal	8
Contraction Expansion	8
Stop word Removal	8
Lemmatization	9
2.4 Visualization	9
3. METHODOLOGY	11
3.1 Tokenization and Padding	11
3.2 Model Architectures	11
3.3 Loss Function	12
3.4 Optimizer	12
3.5 Hyperparameters	12
4. MODEL ARCHITECTURES	13
4.1 Model 1 – Simple RNN with Trainable Embedding	13
4.2 Model 2 – LSTM with Trainable Embedding	13
4.3 Model 3 – LSTM with Pre-trained GloVe Embeddings	14
5. TRAINING CONFIGURATION	14
5.1 Loss Function	14
5.2 Optimizer	14
5.3 Evaluation Metrics	15
5.4 Class Weights	15
5.5 Callbacks	15
6. EXPERIMENTAL ANALYSIS AND RESULTS	15
6.1 Performance Comparison Between RNN and LSTM Models	15
6.2 Computational Efficiency	18

6.3 Training with Different Embeddings	18
6.4 Model Evaluation.....	20
Model 1 — Simple RNN (Test Accuracy: 56.64%)	20
Model 2 — LSTM (Test Accuracy: 79.46%)	20
Model 3 — LSTM + GloVe (Test Accuracy: 82.49%)	20
7. ANALYSIS OF PREDICTION ERRORS	21
7.1 Misclassified Examples	21
7.2 Root Causes	21
Short Sequences and Insufficient Context	21
Category Vocabulary Overlap	21
Class Imbalance	21
Vanishing Gradients in SimpleRNN.....	22
Out-of-Vocabulary Entities.....	22
7.3 Suggested Improvements	22
8. CROSS-MODEL COMPARISON	22
9. GRAPHICAL USER INTERFACE FOR LIVE TEXT CLASSIFICATION.....	25
9.1 Interface Features.....	25
9.2 Prediction Pipeline	25
9.3 Example Predictions	25
10. CONCLUSION AND FUTURE WORK	26
10.1 Key Findings.....	26
Pre-trained Embeddings Are the Most Impactful Lever.....	26
LSTM Substantially Outperforms SimpleRNN.....	26
Class Imbalance Remains a Persistent Challenge.....	26
Preprocessing Quality Matters.....	26
10.2 Limitations	26
10.3 Future Work	27
REFERENCES	27

Table of Figure

Figure 1: Category Distribution — SPORTS dominates at ~41%, creating a 5x class imbalance.....	7
Figure 2: Word Cloud — All Headlines (Cleaned). Dominant terms include "game", "win", "player" (SPORTS), "stock", "rate" (MONEY), and "student", "school" (EDUCATION).....	9
Figure 3: Top 20 Most Frequent Words After Cleaning. Sports-related vocabulary dominates due to class imbalance.	10

Figure 4: Word Clouds by Category. Each panel displays the most distinctive vocabulary for that class. CULTURE & ARTS and EDUCATION share considerable vocabulary overlap..	10
Figure 5: Sequence Length Distribution (Training Set). The red dashed line marks MAX_LEN=12 at the 95th percentile. Mean length is 6.8 tokens, confirming the short-text nature of the task.	11
Figure 6: Simple RNN — Training and Validation Loss / Accuracy Curves. Validation loss begins to diverge around epoch 10, indicating mild overfitting. Validation accuracy plateaus below 60%	16
Figure 7: Simple RNN Confusion Matrix. The model overwhelmingly predicts SPORTS for uncertain inputs. TECH classification is near-zero ($F1 \approx 0.00$), with most TECH headlines misclassified as SPORTS.....	17
Figure 8: LSTM — Training and Validation Loss / Accuracy Curves. Validation accuracy rises steadily to ~79%, demonstrating significantly more stable convergence than Simple RNN	17
Figure 9: LSTM Confusion Matrix. All five classes are now correctly identified with reasonable recall. CULTURE & ARTS and EDUCATION remain the most confused pair due to shared vocabulary.....	18
Figure 10: LSTM + GloVe — Training and Validation Loss / Accuracy Curves. The model plateaus rapidly around epoch 10, benefiting from pre-trained semantic initialization. Validation accuracy reaches ~82%	19
Figure 11: LSTM + GloVe Confusion Matrix. Best overall distribution across all five classes. SPORTS precision remains high (~0.95). CULTURE & ARTS and EDUCATION confusion is reduced compared to Models 1 and 2	19
Figure 12: Overlaid Training and Validation Curves — All Models. SimpleRNN (dark) shows early divergence; LSTM (medium) converges more stably; LSTM + GloVe (light) plateaus fastest at the highest validation accuracy.....	23
Figure 13: Test Accuracy Comparison Bar Chart. The GloVe-enhanced LSTM achieves 82.49%, compared to 79.46% for LSTM and 56.64% for SimpleRNN — a 25.85 percentage-point improvement from Model 1 to Model 3	23
Figure 14: Side-by-Side Confusion Matrix Comparison. Left: SimpleRNN shows heavy SPORTS bias and TECH collapse. Centre: LSTM recovers all classes. Right: LSTM + GloVe achieves the most balanced distribution across all five categories.....	24
Figure 15: Per-Class F1 Score Comparison Across All Three Models. The most dramatic improvements from Model 1 to Model 3 occur in TECH ($0.00 \rightarrow 0.82$) and EDUCATION ($0.25 \rightarrow 0.68$), demonstrating the utility of pre-trained representations for minority classes.	24

1. INTRODUCTION

Automatic text categorization has become a fundamental research area within Natural Language Processing (NLP) due to the explosive growth of digital content available on modern online platforms. News headline classification — the task of automatically assigning short text snippets to predefined topic labels — underpins many practical applications including personalized content delivery, automated news aggregation, misinformation flagging, and trend analysis. Large-scale outlets such as Google News, BBC, and Reuters process millions of articles daily, making manual categorization economically infeasible at that volume.

Historically, text classification systems relied on traditional machine learning algorithms such as Naive Bayes and Support Vector Machines (SVM). These approaches require manual feature engineering to convert text into numerical representations, and while effective for simpler tasks, they fail to capture the sequential and contextual relationships inherent in natural language. Bag-of-words representations, for instance, discard word order entirely — a critical limitation for short texts where every token matters.

The advent of deep learning introduced recurrent neural network architectures that process text as ordered sequences. Standard Recurrent Neural Networks (RNNs) can, in principle, capture temporal dependencies. However, they suffer from the vanishing gradient problem, where gradient signals diminish during backpropagation through time, causing the model to forget information from earlier tokens even in sequences as short as seven words. Long Short-Term Memory (LSTM) networks address this through a gated cell-state mechanism, allowing selective retention of relevant information across longer contexts.

A complementary approach to architecture design is the choice of word representation. Randomly initialized trainable embeddings must be learned from scratch using only the available training data, which can be insufficient for rare or domain-specific vocabulary. Pre-trained word embeddings such as GloVe (Global Vectors for Word Representation), trained on billions of tokens from Wikipedia and Giga word, encode rich semantic relationships. Initializing the embedding layer with these vectors provides the model with linguistic prior knowledge before any task-specific training begins.

This project investigates and compares three recurrent neural network architectures for five-class news headline classification, systematically examining the contribution of both recurrent architecture type and embedding strategy to overall performance:

- Model 1 — Simple RNN with trainable embedding layer
- Model 2 — LSTM with trainable embedding layer
- Model 3 — LSTM with pre-trained and frozen GloVe-50 embeddings

The dataset comprises 11,025 news headlines across five categories: SPORTS, TECH, MONEY, CULTURE & ARTS, and EDUCATION. The project covers the complete NLP pipeline from raw text cleaning through model training and evaluation, culminating in a live prediction interface built with Grado. The primary finding is that embedding quality contributes more to performance gains than architectural complexity for this short-text, small-dataset classification task.

2. DATASET

2.1 Dataset Source

The dataset used in this project is a curated subset of the HuffPost News Category Dataset, originally published on Kaggle by Rishabh Misra in 2022. The complete dataset contains over 200,000 news articles published by HuffPost between 2012 and 2022, spanning 42 topic categories. For the 6CS012 NLP Portfolio coursework (Variant 4 Slate), this was reduced to five categories — SPORTS, TECH, MONEY, CULTURE & ARTS, and EDUCATION — and filtered to include only the headline column, excluding full article text. This deliberate restriction to headlines makes the task more challenging because each input contains minimal contextual tokens.

2.2 Dataset Characteristics

The dataset contains the following key properties:

Property	Value
Total headlines	11,025
Number of classes	5
Vocabulary size (raw)	11,998 unique tokens
Vocabulary size (capped)	5,000 (top-frequency words)
Train / Test split	8,820 / 2,205 (80/20, stratified)
Average tokens per headline	~6.8 (post-cleaning)
95th-percentile sequence length	12 tokens
Imbalance ratio	5.01x (SPORTS vs. EDUCATION)




Figure 1: Category Distribution — SPORTS dominates at ~41%, creating a 5x class imbalance.

The dataset exhibits significant class imbalance. SPORTS is the dominant class with approximately 5,077 headlines (~41%), while EDUCATION is the smallest with ~1,014 headlines (~8%). This roughly 5:1 ratio introduces a strong class bias during training. To counteract this, class weights were computed and applied during training using sklearn's `compute_class_weight` function:

Category	Count	Class Weight	Proportion
SPORTS	5,077	0.434	~41%
TECH	2,104	1.048	~16%
MONEY	1,756	1.256	~14%
CULTURE & ARTS	1,074	2.054	~9%
EDUCATION	1,014	2.175	~8%

The tokenizer was fitted exclusively on training data to prevent any information leakage from the test set into model training. All 2,205 test samples were held out from all preprocessing statistics.

2.3 Data Preprocessing

Prior to tokenization, all headlines were subjected to a systematic text cleaning pipeline to normalize the raw text and reduce noise. The following steps were applied sequentially:

Lowercasing

All text was converted to lowercase to ensure consistent token representations. For example:

"Breaking NEWS on Apple Technology" → "breaking news on apple technology"

URL and Mention Removal

Web addresses beginning with `http` or `https`, Twitter-style `@mentions`, and `#hashtag` tokens were stripped using regular expressions, as they carry no semantic content relevant to category prediction.

Number and Special Character Removal

Digits and punctuation marks were removed. For example, "U.S." becomes "US" and numeric strings such as "2024" are removed entirely. Underscores were replaced with spaces.

Contraction Expansion

Shortened forms were expanded to their full forms using the `contractions` library:

- "don't" → "do not"
- "can't" → "cannot"
- "it's" → "it is"

Stop word Removal.

High-frequency, low-information English words from NLTK's stop word corpus were removed (e.g., "the", "is", "at", "and"). Although stop words carry grammatical function, they contribute minimal discriminative signal for topic classification.

Lemmatization

Words were reduced to their canonical base form using NLTK's WordNet Lemmatize to unify morphological variants:

- "running" → "run"
- "studies" → "study"
- "better" → "good"

After this pipeline, the average headline length dropped to approximately 6.8 tokens. This is notably short, meaning a single ambiguous or out-of-vocabulary word can significantly shift the prediction. The tokenizer was fitted on the cleaned training data with a vocabulary cap of 5,000 words; any word outside this cap was mapped to the <OOV> token.

2.4 Visualization

Following preprocessing, several visualizations were produced to understand the structure and content of the cleaned dataset.



Figure 2: Word Cloud — All Headlines (Cleaned). Dominant terms include "game", "win", "player" (SPORTS), "stock", "rate" (MONEY), and "student", "school" (EDUCATION)




Figure 3: Top 20 Most Frequent Words After Cleaning. Sports-related vocabulary dominates due to class imbalance.

The word cloud confirms the dominance of sports-related vocabulary across all headlines, consistent with SPORTS representing 41% of the dataset. Terms such as "game", "win", "season", and "player" appear largest. Financial terms ("stock", "rate", "fund") and educational terms ("school", "student", "program") are clearly distinct, while CULTURE & ARTS and EDUCATION share overlapping academic vocabulary — a known classification challenge.

Word Clouds by Category



Figure 4: Word Clouds by Category. Each panel displays the most distinctive vocabulary for that class. CULTURE & ARTS and EDUCATION share considerable vocabulary overlap..

The per-category word clouds reveal vocabulary patterns that will directly influence model performance. SPORTS shows strong action verbs and proper nouns (team and player names). MONEY is characterized by financial terminology. Critically, CULTURE & ARTS and EDUCATION share

overlapping vocabulary (study, award, program, research) explaining the persistent confusion between these classes across all models.

3. METHODOLOGY

3.1 Tokenization and Padding

After text cleaning, the preprocessed headlines were transformed into numerical sequences using Keres’s Tokenizer with a vocabulary cap of 5,000 words. Each unique word within the vocabulary was assigned a unique integer index. Words encountered during test-set processing that were absent from the training vocabulary were mapped to the reserved <OOV> (out-of-vocabulary) token.

Example tokenization:

- Cleaned sentence: "deep learning improves accuracy."
- Numerical sequence: [15, 27, 48, 102]

Because neural network models require fixed-length input tensors, all sequences were padded or truncated to a uniform length of 12 tokens. This value was chosen as the approximate 95th percentile of training-set sequence lengths, balancing coverage against unnecessary padding waste. Short sequences were padded with zeros at the end (post-padding); sequences exceeding 12 tokens were truncated from the end (post-truncating).




Figure 5: Sequence Length Distribution (Training Set). The red dashed line marks MAX_LEN=12 at the 95th percentile. Mean length is 6.8 tokens, confirming the short-text nature of the task.

The distribution confirms that most headlines fall well below the 12-token limit, meaning most sequences are padded rather than truncated. The resulting padded arrays have shape (8820, 12) for training and (2205, 12) for testing.

3.2 Model Architectures

Three sequential deep learning architectures were implemented using TensorFlow/Keras:

#	Model	Embedding	Recurrent Layer
1	Simple RNN	Trainable (64-dim)	Simple RNN (32 units)
2	LSTM	Trainable (64-dim)	LSTM (32 units)
3	LSTM + GloVe	Pre-trained GloVe-50 (frozen)	LSTM (32 units)

3.3 Loss Function

Sparse categorical cross-entropy was used for all three models. This is the standard loss function for multi-class classification with integer class labels. It computes the log-probability of the correct class from the model's SoftMax output distribution, penalizing confident incorrect predictions more heavily. Binary cross-entropy would be inappropriate here given five mutually exclusive categories.

3.4 Optimizer

The Adam optimizer was employed with an initial learning rate of 0.001. Adam's adaptive per-parameter learning rates make it robust to the sparse word-frequency distributions typical of text data. A Reduce LR On Plateau callback was configured to halve the learning rate whenever validation loss failed to improve for 2 consecutive epochs (minimum learning rate: 1e-5). An Early Stopping callback with patience=5 restored the best model weights and halted training when validation loss plateaued.

3.5 Hyperparameters

Hyperparameter	Value	Rationale
Vocabulary Size	5,000	Covers top-frequency tokens; reduces noise
Maximum Sequence Length	12	95th percentile of training lengths
Embedding Dimension (Models 1 & 2)	64	Sufficient capacity for 5,000-word vocabulary
Embedding Dimension (Model 3)	50	Fixed by GloVe vector size
RNN / LSTM Units	32	Lightweight; sufficient for short sequences
Dense Units	32	Intermediate projection before SoftMax
Dropout Rate	0.5	Regularize on small dataset
L2 Regularization	0.001	Applied to Dense layer weights
Recurrent Dropout	0.3	Regularize recurrent connections
Batch Size	32	Compromise between speed and gradient stability

Maximum Epochs	30	Upper bound; early stopping prevents over-training
Optimizer	Adam	Adaptive learning rate; robust to sparse features
Loss Function	Sparse Categorical Cross-Entropy	Standard for integer-labelled multi-class tasks
Train / Test Split	80 / 20 (stratified)	Maintains class proportions in both splits

4. MODEL ARCHITECTURES

4.1 Model 1 – Simple RNN with Trainable Embedding

Architecture stack:

Input (12) → Embedding (5000×64) → Simple RNN(32, recurrent dropout=0.3) → Dropout(0.5) → Dense(32, ReLU, L2=0.001) → Dropout(0.5) → Dense(5, SoftMax)

The Simple RNN processes each token in sequence, combining the current input embedding with the previous hidden state to produce an updated hidden state. This recurrent computation allows the model to, in theory, propagate information across the sequence. However, the vanishing gradient problem means that during backpropagation, gradients diminish exponentially as they flow backward through timesteps. Even across sequences of 6–7 tokens, the first tokens contribute negligible signal to the final prediction.

Key properties:

- Fastest training due to a single weight matrix per timestep
- Lowest parameters count among the three models.
- Baseline architecture to establish minimum performance.
- Susceptible to forgetting early sequence tokens due to gradient vanishing.

4.2 Model 2 – LSTM with Trainable Embedding

Architecture stack:

Input (12) → Embedding (5000×64) → LSTM (32, recurrent dropout=0.3) → Dropout(0.5) → Dense(32, ReLU, L2=0.001) → Dropout(0.5) → Dense(5, Softmax)

The LSTM extends the simple RNN with a gated cell-state mechanism. Three learnable gates control information flow at each timestep:

- Forget gate — determines which information to discard from the cell state.
- Input gate — decides what new information to write into the cell state.
- Output gate — controls which part of the cell state to expose as the hidden state.

These gates allow the LSTM to selectively retain or discard information, substantially mitigating the vanishing gradient problem. Although each timestep requires four matrix multiplications (versus one for Simple RNN), the gating provides far superior long-range information retention. The LSTM uses the same embedding dimension, dropout configuration, and dense head as Model 1, isolating the recurrent layer as the only architectural variable between Models 1 and 2.

4.3 Model 3 – LSTM with Pre-trained GloVe Embeddings

Architecture stack:

Input (12) → Embedding (5000 × 50, GloVe-frozen) → LSTM (32, recurrent dropout=0.3) → Dropout(0.5) → Dense(32, ReLU, L2=0.001) → Dropout(0.5) → Dense(5, Softmax)

Model 3 uses the identical LSTM architecture as Model 2 but replaces the randomly initialized embedding layer with a pre-trained GloVe-50 embedding matrix (glove-wiki-gigaword-50, loaded via the genism downloader API). GloVe vectors were trained on 6 billion tokens from Wikipedia and Giga word corpora across a 400,000-word vocabulary.

An embedding matrix of shape (5000 × 50) was constructed by looking up each word in the training vocabulary within the GloVe index. Words found in GloVe received their pre-trained 50-dimensional vector; words not found (including <OOV>) were assigned zero vectors. The embedding layer was set to trainable=False, freezing these weights throughout training. This prevents the model from degrading the pre-learned semantic relationships and reduces the total number of learnable parameters.

Benefits of GloVe initialization:

- Semantic proximity encoded before any task-specific training (e.g., "fund", "stock", "market" cluster together)
- Fewer learnable parameters — the embedding layer does not backpropagate
- Improved generalization on small datasets where random embeddings underfit
- Faster convergence as meaningful word representations are immediately available.

Word coverage analysis showed that a significant proportion of the capped 5,000-word vocabulary was found in the GloVe index, providing rich initialization for the most common news headline vocabulary.

5. TRAINING CONFIGURATION

5.1 Loss Function

Sparse categorical cross-entropy was selected as the optimization objective because this is a multi-class classification task with integer labels (0 through 4 corresponding to the five categories). This loss function computes the negative log-likelihood of the true class under the predicted Softmax probability distribution, providing stronger gradient signals for confidently incorrect predictions and stable learning across all five output classes.

5.2 Optimizer

The Adam optimizer (Adaptive Moment Estimation) was chosen for its robustness on text data, where word frequencies are highly skewed. Adam maintains per-parameter first and second moment estimates, effectively providing a per-parameter adaptive learning rate. Compared to standard SGD, Adam converges more reliably on sparse gradient landscapes typical of high-vocabulary NLP tasks.

A Reduce LR On Plateau callback monitored validation loss and halved the learning rate whenever no improvement was observed for 2 consecutive epochs. This progressive reduction prevents overshooting in later epochs and helps fine-tune the model toward a lower loss minimum. Early Stopping with `patience=5` and `restore_best_weights = True` ensured the final saved weights corresponded to the epoch with the lowest validation loss.

5.3 Evaluation Metrics

The following metrics were used to comprehensively evaluate model performance:

- Classification Accuracy — proportion of correctly classified test headlines
- Precision — for each class, the fraction of predicted positives that are genuine positives.
- Recall — for each class, the fraction of actual positives correctly identified.
- F1-Score — harmonic mean of precision and recall; preferred over accuracy under class imbalance.
- Macro-average F1 — unweighted average F1 across all five classes, penalizing class-specific failures equally
- Confusion Matrix — visualizes per-class prediction patterns and common misclassification pairs.

5.4 Class Weights

To counteract the 5:1 class imbalance between SPORTS and EDUCATION, class weights were computed using sklearn's `compute_class_weight("balanced")` function and passed to `model.fit()` via the `class_weight` argument. Higher weights assigned to minority classes (EDUCATION: 2.175, CULTURE & ARTS: 2.054) penalize misclassification of underrepresented categories more heavily during gradient updates, encouraging the model to learn from minority-class examples rather than defaulting to the dominant SPORTS class.

5.5 Callbacks

Each model training run was equipped with fresh callback instances to prevent state carry-over between experiments:

- Early Stopping: `monitor=val_loss`, `patience=5`, `restore_best_weights=True` — halts training when the model stops improving and resets to the best checkpoint
- ReduceLROnPlateau: `monitor=val_loss`, `factor=0.5`, `patience=2`, `min_lr=1e-5` — progressively reduces the learning rate to enable finer optimization in later epochs

6. EXPERIMENTAL ANALYSIS AND RESULTS

6.1 Performance Comparison Between RNN and LSTM Models

Model	Test Accuracy	Macro F1	Behavior
Simple RNN	56.64%	~0.37	Diverged ~epoch 10
LSTM	79.46%	~0.73	Stable to ~epoch 15
LSTM + GloVe	82.49%	~0.78	Plateaued ~epoch 10

The Simple RNN achieved a test accuracy of 56.64% — well above the majority-class baseline of ~41% (always predicting SPORTS), but substantially below the LSTM models. Notably, the Simple RNN's TECH F1 was approximately zero, with the model collapsing to the SPORTS majority class for uncertain inputs. Despite the short average sequence length of 6.8 tokens, the vanishing gradient problem prevents Simple RNN from reliably propagating discriminative signal from earlier tokens.

The LSTM improved accuracy dramatically to 79.46%, recovering class-specific performance across all five categories. The gating mechanism proved beneficial even at these short sequence lengths. LSTM training was approximately 3x slower per epoch due to four gate matrix operations per timestep but converged more stably with smaller validation loss oscillation.

Training Curves — Model 1: SimpleRNN




Figure 6: Simple RNN — Training and Validation Loss / Accuracy Curves. Validation loss begins to diverge around epoch 10, indicating mild overfitting. Validation accuracy plateaus below 60%




Figure 7: Simple RNN Confusion Matrix. The model overwhelmingly predicts *SPORTS* for uncertain inputs. *TECH* classification is near-zero ($F1 \approx 0.00$), with most *TECH* headlines misclassified as *SPORTS*.




Figure 8: LSTM — Training and Validation Loss / Accuracy Curves. Validation accuracy rises steadily to ~79%, demonstrating significantly more stable convergence than Simple RNN




Figure: 9LSTM Confusion Matrix. All five classes are now correctly identified with reasonable recall. CULTURE & ARTS and EDUCATION remain the most confused pair due to shared vocabulary.

6.2 Computational Efficiency

Model	Relative Speed	Key Cost	Notes
Simple RNN	Fastest	1 weight matrix/step	Lowest memory usage
LSTM	~3x slower	4 gate matrices/step	Better generalization
LSTM + GloVe	Like LSTM	GloVe setup (one-off)	Frozen weights speed up backprop

All experiments were conducted on Google Collab with a T4 GPU (8GB VRAM). The Simple RNN was the fastest per epoch due to its minimal weight matrix. Although LSTM incurs four matrix operations per timestep, the frozen GloVe embedding layer in Model 3 reduces backpropagation overhead, resulting in similar epoch times to LSTM Model 2. The one-off GloVe download and embedding matrix construction added a small upfront cost but is not repeated across training epochs.

6.3 Training with Different Embeddings

Training Curves — Model 3: LSTM + Word2Vec




Figure 10: LSTM + GloVe — Training and Validation Loss / Accuracy Curves. The model plateaus rapidly around epoch 10, benefiting from pre-trained semantic initialization. Validation accuracy reaches ~82%




Figure 11: LSTM + GloVe Confusion Matrix. Best overall distribution across all five classes. SPORTS precision remains high (~0.95). CULTURE & ARTS and EDUCATION confusion is reduced compared to Models 1 and 2

The comparison between random trainable embeddings (Models 1 and 2) and frozen GloVe embeddings (Model 3) is the central empirical finding of this project. Models 1 and 2 must learn word representations purely from 8,820 short training headlines — an insufficient corpus for reliable semantic learning, particularly for rare words and named entities. GloVe vectors, trained on 6 billion tokens, provide immediate semantic proximity between related terms (e.g., "stock", "fund", "market" are already clustered; "player", "team", "score" are grouped) before task-specific training begins.

This initialization advantage translates to a 13-percentage-point accuracy improvement (LSTM: 79.46% → LSTM + GloVe: 82.49%) and a 0.05 macro F1 gain, without any change to the LSTM architecture itself. This result demonstrates that for small-dataset, short-text classification, embedding quality is the primary performance lever — more impactful than architectural complexity.

6.4 Model Evaluation

Classification Reports:

Model 1 — Simple RNN (Test Accuracy: 56.64%)

Class	Precision	Recall	F1-Score	Support
CULTURE & ARTS	0.18	0.26	0.21	215
EDUCATION	0.20	0.34	0.25	203
MONEY	0.36	0.59	0.45	351
SPORTS	0.93	0.90	0.92	1015
TECH	0.00	0.00	0.00	421
Macro Average	0.34	0.42	0.37	2205

Model 2 — LSTM (Test Accuracy: 79.46%)

Class	Precision	Recall	F1-Score	Support
CULTURE & ARTS	0.66	0.80	0.72	215
EDUCATION	0.49	0.63	0.55	203
MONEY	0.68	0.76	0.71	351
SPORTS	0.94	0.89	0.91	1015
TECH	0.86	0.68	0.76	421
Macro Average	0.73	0.75	0.73	2205

Model 3 — LSTM + GloVe (Test Accuracy: 82.49%)

Class	Precision	Recall	F1-Score	Support
CULTURE & ARTS	0.63	0.82	0.71	215

EDUCATION	0.62	0.74	0.68	203
MONEY	0.79	0.75	0.77	351
SPORTS	0.95	0.87	0.91	1015
TECH	0.83	0.81	0.82	421
Macro Average	0.76	0.80	0.78	2205

7. ANALYSIS OF PREDICTION ERRORS

7.1 Misclassified Examples

A qualitative analysis of misclassified test headlines reveals patterns in model failure modes. The following examples illustrate common error types across all three models:

Headline	True Label	Predicted	Error Reason
Apple will reduce the iPhone price before holidays	MONEY	TECH (RNN)	Brand name "Apple" biases toward TECH
New scholarship fund launched at Harvard	EDUCATION	MONEY (RNN)	"fund" is a strong MONEY signal
Grammy ceremony celebrates best in music and culture	CULTURE & ARTS	EDUCATION (LSTM)	Shared vocabulary: award, celebrate
Player contract deal sets record valuation	SPORTS	MONEY (LSTM+GloVe)	Financial terms override sports context

7.2 Root Causes

Short Sequences and Insufficient Context

With an average of only 6.8 cleaned tokens per headline, a single ambiguous word can dominate the prediction. For example, "fund" in an education headline strongly activates MONEY class neurons, overriding the broader educational context that would be apparent in longer text.

Category Vocabulary Overlap

CULTURE & ARTS and EDUCATION share substantial vocabulary: words such as "study", "award", "school", "research", and "program" appear frequently in both categories. This lexical ambiguity is difficult for any model to resolve from the headline alone and likely requires more context than a single short sentence provides. SPORTS and MONEY also overlap on athlete contracts, valuations, and sponsorship headlines.

Class Imbalance

Despite class weighting, the 5:1 SPORTS-to-EDUCATION ratio creates a strong prior bias. The Simple RNN completely fails to classify TECH headlines ($F1 = 0.00$) because defaulting to SPORTS is optimally penalized less by the loss function before class weights take full effect.

Vanishing Gradients in SimpleRNN

Even at 6–7 tokens, the SimpleRNN fails to propagate sufficient gradient from the final classification back to the first token position. The LSTM's cell-state mechanism resolves this, which is why LSTM recovers MONEY F1 from effectively zero (SimpleRNN) to 0.71.

Out-of-Vocabulary Entities

Named entities — player names, company names, and rare proper nouns — frequently appear fewer than five times in 8,820 training headlines and fall outside the 5,000-token vocabulary cap. These collapse to the <OOV> token, removing potentially strong discriminative signals.

7.3 Suggested Improvements

- Bidirectional LSTM — reads sequences in both directions, providing each token with access to both preceding and following context, doubling effective receptive field.
- Fine-tuned GloVe — unfreeze the embedding layer after a frozen warm-up phase, allowing domain adaptation to news vocabulary.
- Larger vocabulary — extend from 5,000 to 15,000–20,000 words to reduce OOV entity losses.
- Data augmentation — apply synonym substitution or back-translation to minority classes (CULTURE & ARTS, EDUCATION) to reduce the imbalance.
- Domain-specific embeddings — use GloVe-Twitter or a news-corpus-trained model rather than Wikipedia/Gigaword for better headline domain coverage.
- Transformer-based models — DistilBERT or RoBERTa fine-tuned on news headlines would likely exceed 90% accuracy through contextual bidirectional attention.
- Hyperparameter search — systematic grid or Bayesian optimization for LSTM units, dropout rate, embedding dimensions, and learning rate.

8. CROSS-MODEL COMPARISON

The following visualizations provide a side-by-side comparison of all three model architectures, enabling direct assessment of performance differences across training behavior, final accuracy, classification patterns, and per-class F1 scores.




Figure 12: Overlaid Training and Validation Curves — All Models. SimpleRNN (dark) shows early divergence; LSTM (medium) converges more stably; LSTM + GloVe (light) plateaus fastest at the highest validation accuracy.

The overlaid curves confirm that LSTM + GloVe achieves the fastest and most stable convergence — plateau reached around epoch 10 with minimal validation loss oscillation. The LSTM without GloVe continues improving until ~epoch 15 before early stopping. The SimpleRNN exhibits the characteristic divergence pattern where training loss decreases while validation loss stagnates or increases, indicating model capacity wasted on training-set memorization rather than generalization.




Figure 13: Test Accuracy Comparison Bar Chart. The GloVe-enhanced LSTM achieves 82.49%, compared to 79.46% for LSTM and 56.64% for SimpleRNN — a 25.85 percentage-point improvement from Model 1 to Model 3




Figure 14: Side-by-Side Confusion Matrix Comparison. Left: SimpleRNN shows heavy SPORTS bias and TECH collapse. Centre: LSTM recovers all classes. Right: LSTM + GloVe achieves the most balanced distribution across all five categories

The confusion matrices reveal the qualitative improvement trajectory across models. SimpleRNN's matrix shows strong SPORTS prediction dominance with most non-SPORTS headlines misclassified into SPORTS or MONEY. LSTM's matrix shows successful discrimination across all classes, with the primary remaining confusion between CULTURE & ARTS and EDUCATION. LSTM + GloVe shows the most balanced diagonal, confirming that pre-trained embeddings improve both accuracy and class equity.




Figure 15: Per-Class F1 Score Comparison Across All Three Models. The most dramatic improvements from Model 1 to Model 3 occur in TECH (0.00 → 0.82) and EDUCATION (0.25 → 0.68), demonstrating the utility of pre-trained representations for minority classes.

The per-class F1 comparison is the most informative view of model progression. TECH is the most striking case: SimpleRNN achieves F1 = 0.00 (complete failure), LSTM achieves 0.76, and LSTM + GloVe achieves 0.82. EDUCATION improves from 0.25 to 0.68. SPORTS remains consistently high across all models (~0.91–0.92), confirming that the dominant class benefits all architectures. CULTURE & ARTS is the only class where LSTM + GloVe (0.71) does not improve upon LSTM (0.72), likely because the two models share the same architecture, and the CULTURE & ARTS / EDUCATION overlap remains a hard challenge even with better embeddings.

9. GRAPHICAL USER INTERFACE FOR LIVE TEXT CLASSIFICATION

A Gradio-based interactive web application was developed to enable real-time headline classification using all three trained models. The interface was deployed using Gradio's `share=True` parameter, generating a publicly accessible temporary URL. The application successfully classified 5 out of 6 pre-loaded example headlines correctly, with the best model achieving confidence scores up to 96%.

9.1 Interface Features

- Free-text input field for typing or pasting any news headline.
- Dropdown model selector to switch between SimpleRNN, LSTM, and LSTM + GloVe pipelines.
- Prediction panel displaying the predicted category, confidence percentage, and cleaned headline.
- Per-class confidence bar chart (`gr.Label`) showing probability distribution across all five categories
- Six pre-loaded example headlines covering all categories, launchable with a single click.
- Soft "slate" theme aesthetic matching the Variant 4 Slate palette specification.

9.2 Prediction Pipeline

The prediction pipeline within the Gradio interface exactly mirrors the training preprocessing to ensure consistency. When a user submits a headline, the following steps execute in sequence:

- Step 1: Apply the same `clean_text()` function (lowercase → expand contractions → remove URLs/mentions → remove digits/punctuation → remove stopwords → lemmatize)
- Step 2: Tokenise the cleaned string using the fitted Keras Tokenizer (vocabulary capped at 5,000 words; OOV mapped to `<OOV>` token)
- Step 3: Pad or truncate the resulting sequence to `MAX_LEN=12` tokens (post-padding)
- Step 4: Pass the padded tensor through the selected model to obtain a softmax probability distribution over five classes.
- Step 5: Return the `argmax` class label and confidence percentage alongside the full per-class distribution.

The use of identical preprocessing in both training and inference prevents train-test skew and ensures that the model receives input in exactly the format it was trained on. The interface demonstrates the practical deployability of the trained classification pipeline beyond the notebook environment.

9.3 Example Predictions

Example Headline	True Category	Model Used	Result
Manchester United wins Premier League title	SPORTS	LSTM	Correct ✓
Apple unveils new AI chip for iPhone	TECH	LSTM	Correct ✓

Federal Reserve raises interest rates amid inflation	MONEY	LSTM + GloVe	Correct ✓
Grammy Awards celebrates best in music and culture	CULTURE & ARTS	SimpleRNN	Correct ✓
Harvard University launches new scholarship program	EDUCATION	LSTM + GloVe	Correct ✓
Tesla stock surges after recording quarterly earnings	MONEY	LSTM	Correct ✓

10. CONCLUSION AND FUTURE WORK

10.1 Key Findings

This project successfully implemented and evaluated three deep learning architectures for automated news headline classification across five categories. The key experimental findings are as follows:

Pre-trained Embeddings Are the Most Impactful Lever

Replacing randomly initialized trainable embeddings with frozen GloVe-50 vectors delivered a 13-percentage-point accuracy improvement (79.46% → 82.49%) and a macro F1 gain of 0.05, without any modification to the LSTM architecture. This confirms that for short-text, small-dataset tasks, representation quality outweighs model complexity.

LSTM Substantially Outperforms SimpleRNN

The LSTM's gating mechanism was effective even across these short 6.8-token sequences, recovering TECH classification from F1=0.00 to F1=0.76 and improving overall accuracy from 56.64% to 79.46%. The vanishing gradient limitation of SimpleRNN manifests even on very short sequences, confirming the practical necessity of gated architectures for NLP classification.

Class Imbalance Remains a Persistent Challenge

Despite balanced class weighting, the 5:1 SPORTS-to-EDUCATION ratio continues to influence per-class behavior. SPORTS achieves consistently high F1 (~0.91–0.92) across all models, while minority classes — particularly EDUCATION and CULTURE & ARTS — benefit most from the GloVe initialization. The vocabulary overlap between CULTURE & ARTS and EDUCATION is the primary remaining error source across all three models.

Preprocessing Quality Matters

The multi-step cleaning pipeline — contraction expansion, lemmatization, stop word removal, and percentile-based padding — provided clean, consistent inputs that contributed to training stability. The tokenizer-on-training-only approach prevented data leakage and ensured unbiased test evaluation.

10.2 Limitations

- Limited to five categories — real-world news classification involves many more classes.
- Short headline inputs (average 6.8 tokens) restrict contextual learning regardless of architecture.
- Frozen GloVe vectors cannot adapt to domain-specific vocabulary (sports team names, tech product names, new financial instruments)
- No hyperparameter optimization — values were selected based on established practice, not systematic search.
- Single train/test split — k-fold cross-validation would provide more reliable performance estimates.

10.3 Future Work

- Bidirectional LSTM / GRU — processes sequences in both directions, particularly beneficial for short headlines
- GloVe fine-tuning — start with frozen embeddings and unfreeze after initial convergence to enable domain adaptation.
- DistilBERT or RoBERTa — state-of-the-art transformer models fine-tuned on news headlines are expected to exceed 90% accuracy.
- Data augmentation — synonym substitution and back-translation for minority classes to reduce imbalance.
- Larger vocabulary (15,000–20,000) — better coverage of named entities and domain-specific terms.
- Ensemble methods — soft-voting across SimpleRNN, LSTM, and LSTM + GloVe may yield marginal accuracy improvements.
- API deployment — wrapping the best model in a REST API using FastAPI or Flask for production integration.

Overall, this project provided comprehensive hands-on experience in building a complete NLP pipeline — from raw text cleaning through model training, evaluation, error analysis, and live deployment — using modern deep learning techniques. The clear progression from SimpleRNN to LSTM to LSTM + GloVe provides a methodologically clean ablation study of both architecture and representation choices for short-text classification.

REFERENCES

- Bird, S., Klein, E., & Loper, E. (2009). *Natural Language Processing with Python*. O'Reilly Media.
- Chollet, F. (2018). *Deep Learning with Python*. Manning Publications.
- Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2018). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. *Proceedings of NAACL 2019*.
- Hochreiter, S., & Schmidhuber, J. (1997). Long Short-Term Memory. *Neural Computation*, 9(8), 1735–1780.

- Kim, Y. (2014). Convolutional Neural Networks for Sentence Classification. Proceedings of EMNLP 2014.
- Misra, R. (2022). News Category Dataset. Kaggle. <https://www.kaggle.com/datasets/rmisra/news-category-dataset>
- Pennington, J., Socher, R., & Manning, C. D. (2014). GloVe: Global Vectors for Word Representation. Proceedings of EMNLP 2014.
- Pedregosa, F., et al. (2011). Scikit-learn: Machine Learning in Python. Journal of Machine Learning Research, 12, 2825–2830.
- Řehůřek, R., & Sojka, P. (2010). Software Framework for Topic Modelling with Large Corpora. LREC 2010 Workshop.
- Abadi, M., et al. (2015). TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems. Available at: tensorflow.org
- Zhang, X., Zhao, J., & LeCun, Y. (2015). Character-level Convolutional Networks for Text Classification. NeurIPS 2015.